

Screenshots

Execution steps:

```
Mangalas-MBP:WordCountV3 mangalakhandekar$ cd ..
Mangalas-MBP:HadoopPrograms mangalakhandekar$ cd WordCountV4
Mangalas-MBP:WordCountV4 mangalakhandekar$ javac WordCountV4.java -cp $(hadoop classpath)
Mangalas-MBP:WordCountV4 mangalakhandekar$ jar cf WordCountV4.jar WordCount*.class
Mangalas-MBP:WordCountV4 mangalakhandekar$ time hadoop jar WordCountV4.jar WordCountV4 /user/mangalakhandekar/inputLarge /user/mangalakhandekar/outputLarge/wcV4Output -skip /user/mangalakhandekar/stop_words.txt
2018-09-18 19:44:54,907 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
2018-09-18 19:44:55,745 INFO impl.MetricsConfig: loaded properties from hadoop-metrics2.properties
2018-09-18 19:44:55,875 INFO impl.MetricsSystemImpl: Scheduled Metric snapshot period at 10 second(s).
2018-09-18 19:44:55,875 INFO impl.MetricsSystemImpl: JobTracker metrics system started
2018-09-18 19:44:56,030 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
2018-09-18 19:44:56,157 INFO input.FileInputFormat: Total input files to process : 1
2018-09-18 19:44:56,222 INFO mapreduce.JobSubmitter: number of splits:1
2018-09-18 19:44:56,422 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_local1037996150_0001
2018-09-18 19:44:56,424 INFO mapreduce.JobSubmitter: Executing with tokens: []
2018-09-18 19:44:56,713 INFO mapred.LocalDistributedCacheManager: Creating symlink: /usr/local/Cellar/hadoop/hdfs/tmp/mapred/local/1537314296504/stop_words.txt <- /Users/mangalakhandekar/HadoopPrograms/WordCountV4/stop_words.txt
2018-09-18 19:44:56,720 INFO mapred.LocalDistributedCacheManager: Localized hdfs://localhost:9000/user/mangalakhandekar/stop_words.txt as file:/usr/local/Cellar/hadoop/hdfs/tmp/mapred/local/1537314296504/stop_words.txt
2018-09-18 19:44:56,796 INFO mapreduce.Job: The url to track the job: http://localhost:8080/
2018-09-18 19:44:56,797 INFO mapreduce.Job: Running job: job_local1037996150_0001
2018-09-18 19:44:56,798 INFO mapred.LocalJobRunner: OutputCommitter set in config null
2018-09-18 19:44:56,804 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2018-09-18 19:44:56,804 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup _temporary folders under output directory:false, ignore cleanup failures: false
2018-09-18 19:44:56,805 INFO mapred.LocalJobRunner: OutputCommitter is org.apache.hadoop.mapreduce.lib.output.FileOutputCommitter
2018-09-18 19:44:56,856 INFO mapred.LocalJobRunner: Waiting for map tasks
2018-09-18 19:44:56,857 INFO mapred.LocalJobRunner: Starting task: attempt_local1037996150_0001_m_000000_0
2018-09-18 19:44:56,879 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2018-09-18 19:44:56,879 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup _temporary folders under output directory:false, ignore cleanup failures: false
2018-09-18 19:44:56,886 INFO util.ProcfsBasedProcessTree: ProcfsBasedProcessTree currently is supported only on Linux.
2018-09-18 19:44:56,887 INFO mapred.Task: Using ResourceCalculatorProcessTree : null
2018-09-18 19:44:56,891 INFO mapred.MapTask: Processing split: hdfs://localhost:9000/user/mangalakhandekar/inputLarge/largeFile1.txt:0+408778
2018-09-18 19:44:56,962 INFO mapred.MapTask: (EQUATOR) 0 kvi 26214396(104857584)
2018-09-18 19:44:56,962 INFO mapred.MapTask: mapreduce.task.io.sort.mb: 100
2018-09-18 19:44:56,962 INFO mapred.MapTask: soft limit at 83886080
2018-09-18 19:44:56,962 INFO mapred.MapTask: bufstart = 0; bufvoid = 104857600
2018-09-18 19:44:56,962 INFO mapred.MapTask: kvstart = 26214396; length = 6553600
2018-09-18 19:44:56,966 INFO mapred.MapTask: Map output collector class = org.apache.hadoop.mapred.MapTask$MapOutputBuffer
2018-09-18 19:44:57,129 INFO mapred.LocalJobRunner:
2018-09-18 19:44:57,129 INFO mapred.MapTask: Starting flush of map output
2018-09-18 19:44:57,129 INFO mapred.MapTask: Spilling map output
2018-09-18 19:44:57,129 INFO mapred.MapTask: bufstart = 0; bufend = 581669; bufvoid = 104857600
2018-09-18 19:44:57,129 INFO mapred.MapTask: kvstart = 26214396(104857584); kvend = 25984740(103938960); length = 229657/6553600
2018-09-18 19:44:57,274 INFO mapred.MapTask: Finished spill 0
2018-09-18 19:44:57,292 INFO mapred.Task: Task:attempt_local1037996150_0001_m_000000_0 is done. And is in the process of committing
2018-09-18 19:44:57,295 INFO mapred.LocalJobRunner: map
2018-09-18 19:44:57,295 INFO mapred.Task: Task 'attempt_local1037996150_0001_m_000000_0' done.
2018-09-18 19:44:57,302 INFO mapred.Task: Final Counters for attempt_local1037996150_0001_m_000000_0: Counters: 23
File System Counters
FILE: Number of bytes read=4522
FILE: Number of bytes written=636758
FILE: Number of read operations=0
```

Input file creation on local and copying it to HDFS:

```
Mangalas-MBP:WordCountV1 mangalakhandekar$ hadoop fs -mkdir -p /user/mangalakhandekar/inputSimple
2018-09-18 18:50:36,744 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Mangalas-MBP:WordCountV1 mangalakhandekar$ cd
Mangalas-MBP:~ mangalakhandekar$ hadoop fs -copyFromLocal Desktop/simpleFile1.txt /user/mangalakhandekar/inputSimple
2018-09-18 18:51:14,102 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Mangalas-MBP:~ mangalakhandekar$ hadoop fs -copyFromLocal Desktop/simpleFile2.txt /user/mangalakhandekar/inputSimple
2018-09-18 18:51:22,630 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Mangalas-MBP:~ mangalakhandekar$ hadoop fs -copyFromLocal Desktop/simpleFile3.txt /user/mangalakhandekar/inputSimple
2018-09-18 18:51:32,010 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Mangalas-MBP:~ mangalakhandekar$
```

Output Screenshots:

Below screenshots are samples of output:

Dataset1:

Word Count Program V1:

Bytes Written=100

```
real 0m3.847s
user 0m9.809s
sys 0m0.608s
Mangalas-MBP:WordCountV1 mangalakhandekar$ hadoop fs -cat /user/mangalakhandekar/
outputSimple/wcV1Output/part-r-00000
2018-09-18 19:08:38,294 WARN util.NativeCodeLoader: Unable to load native-hadoop
library for your platform... using builtin-java classes where applicable
! 1
Bye 1
Goodbye 1
Hadoop 2
Hello 2
World 2
and 1
bye 1
goodbye 1
hadoop 2
hello 2
world 2
. 1
Mangalas-MBP:WordCountV1 mangalakhandekar$
```

Word Count Program V2:

Bytes Written=100

```
real 0m3.853s
user 0m9.778s
sys 0m0.615s
Mangalas-MBP:WordCountV2 mangalakhandekar$ hadoop fs -cat /user/mangalakhandekar/
outputSimple/wcV2Output/part-r-00000
2018-09-18 19:10:20,514 WARN util.NativeCodeLoader: Unable to load native-hadoop
library for your platform... using builtin-java classes where applicable
! 1
Bye 1
Goodbye 1
Hadoop 2
Hello 2
World 2
and 1
bye 1
goodbye 1
hadoop 2
hello 2
world 2
. 1
Mangalas-MBP:WordCountV2 mangalakhandekar$
```

Word Count Program V3:

Bytes Written=59

```
real 0m4.092s
user 0m10.139s
sys 0m0.663s
Mangalas-MBP:WordCountV3 mangalakhandekar$ hadoop fs -cat /user/mangalakhandekar/outputSimple/
wcV3Output/part-r-00000
2018-09-18 19:12:17,859 WARN util.NativeCodeLoader: Unable to load native-hadoop library for
your platform... using builtin-java classes where applicable
! 1
and 1
bye 2
goodbye 2
hadoop 4
hello 4
world 4
. 1
```

Word Count Program V4:

Bytes Written=55

```
real 0m4.019s
user 0m9.783s
sys 0m0.659s
Mangalas-MBP:WordCountV4 mangalakhandekar$ hadoop fs -cat /user/mangalakhandekar/
outputSimple/wcV4Output/part-r-00000
2018-09-18 19:15:47,698 WARN util.NativeCodeLoader: Unable to load native-hadoop
library for your platform... using builtin-java classes where applicable
bye 2
goodbye 2
hadoop 4
hello 4
world 4
Mangalas-MBP:WordCountV4 mangalakhandekar$
```

DataSet2:

Word Count Program V1:

Bytes Written=476

```
real0m3.926s
user0m9.720s
sys 0m0.653s
Mangalas-MBP:WordCountV1 mangalakhandekar$ hadoop fs -cat /user/mangalakhandekar/
wcV1Output/part-r-00000
2018-09-18 15:57:37,685 WARN util.NativeCodeLoader: Unable to load native-hadoop
library for your platform... using builtin-java classes where applicable
, 1
. 1
A 2
An 2
And 1
Are 1
Because, 1
But 1
Elephant 1
HDFS 1
Hadoop 2
Hadoop. 1
He 2
Hive, 1
Impala, 1
King! 1
Or 1
Sqoop. 1
The 1
Useful 1
a 1
an 2
and 4
anything 1
at 1
bad, 1
but 1
cling! 1
core 1
```

Word Count Program V2:

Bytes Read=423
File Output Format Counters
Bytes Written=446

```
real0m3.839s
user0m8.994s
sys 0m0.623s
Mangalas-MBP:WordCountV2 mangalakhandekar$ hadoop fs -cat /user/mangalakhandekar/wcV2Output/part-r-00000
2018-09-18 16:00:29,940 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
! 2
, 8
. 8
A 2
An 2
And 1
Are 1
Because 1
But 1
Elephant 1
HDFS 1
Hadoop 3
He 2
Hive 1
Impala 1
King 1
Or 1
Sqoop 1
The 1
Useful 1
a 1
an 2
and 4
anything 1
at 1
bad 1
but 1
cling 1
core 1
data 1
does 1
elegant 2
element 1
elephant 2
extraneous 1
fellow 1
forgets 1
gentle 1
gets 1
```

Word Count Program V3:

```
File Input Format Counters
  Bytes Read=425
File Output Format Counters
  Bytes Written=391
```

```
real 0m3.865s
user 0m8.776s
sys 0m0.631s
Mangalas-MBP:WordCountV3 mangalakhandekar$ hadoop fs -cat /user/mangalakhandekar/wcV3Output/part-r-00000
2018-09-18 16:03:13,870 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
! 2
, 8
. 8
a 3
an 4
and 5
anything 1
are 1
at 1
bad 1
because 1
but 2
cling 1
core1
data1
does1
elegant 2
element 1
elephant 3
extraneous 1
fellow 1
forgets 1
gentle 1
gets1
group 1
hadoop 3
hdfs1
he 3
helps 1
him 1
his 1
hive1
impala 1
in 1
is 5
king2
lets1
mad 1
mellow 1
never 2
or 3
plays 1
sqoop 1
the 4
thing 1
thrive 1
```

Word Count Program V4:

```
File Output Format Counters
  Bytes Written=341
```

```
real 0m3.887s
user 0m9.751s
sys 0m0.625s
Mangalas-MBP:WordCountV4 mangalakhandekar$ hadoop fs -cat /user/mangalakhandekar/wcV4Output/part-r-00000
2018-09-18 18:29:36,556 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
! 2
anything 1
are 1
at 1
bad 1
because 1
cling 1
core1
data1
does1
elegant 2
element 1
elephant 3
extraneous 1
fellow 1
forgets 1
gentle 1
gets1
group 1
hadoop 3
hdfs1
he 3
helps 1
him 1
his 1
hive1
impala 1
in 1
king2
lets1
mad 1
mellow 1
never 2
plays 1
sqoop 1
thing 1
thrive 1
useful 1
well1
what1
with1
wonderful 1
yellow 2
```